

time series

Ayoub Asri

2024-03-26

Contents

R packages	1
data	2
time series packages	2
exploration	4
ts object	5
ARIMA model	8
some ideas	10
Autocorrelation plots	12
ADF test	16
Auto Arima	17
forecast on test	20
comments :	21

R packages

```
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr     1.1.4     v readr     2.1.5
## vforcats   1.0.0     v stringr   1.5.1
## v ggplot2   3.5.0     v tibble    3.2.1
## v lubridate 1.9.3     v tidyv     1.3.1
## v purrr    1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()   masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

data

```
a <- getwd()
file <- paste(str_remove(a,"/time series ideas"), "/data/data_article/hourly_5y_data.csv", sep = "")
data <- read_csv(file)

## Rows: 41607 Columns: 2
## -- Column specification -----
## Delimiter: ","
## dbl (1): measure
## dttm (1): time
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

time series packages

```
library(tseries)

## Registered S3 method overwritten by 'quantmod':
##   method           from
##   as.zoo.data.frame zoo

library(timetk)
library(readxl)
library(tidyquant)

## Loading required package: PerformanceAnalytics

## Loading required package: xts

## Loading required package: zoo

##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##   as.Date, as.Date.numeric

##
## ##### Warning from 'xts' package #####
## #
## # The dplyr lag() function breaks how base R's lag() function is supposed to #
## # work, which breaks lag(my_xts). Calls to lag(my_xts) that you type or      #
## # source() into this session won't work correctly.                            #
## #
## # Use stats::lag() to make sure you're not using dplyr::lag(), or you can add #
```

```

## # conflictRules('dplyr', exclude = 'lag') to your .Rprofile to stop      #
## # dplyr from breaking base R's lag() function.                          #
## #
## # Code in packages is not affected. It's protected by R's namespace mechanism #
## # Set 'options(xts.warn_dplyr_breaks_lag = FALSE)' to suppress this warning. #
## #
## #####
## 
## Attaching package: 'xts'

## The following objects are masked from 'package:dplyr':
## 
##     first, last

## 
## Attaching package: 'PerformanceAnalytics'

## The following object is masked from 'package:graphics':
## 
##     legend

## Loading required package: quantmod

## Loading required package: TTR

library(scales)

## 
## Attaching package: 'scales'

## The following object is masked from 'package:purrr':
## 
##     discard

## The following object is masked from 'package:readr':
## 
##     col_factor

library(forecast)    # forecasting pkg
library(sweep)       # Broom tidiers for forecast pkg
library(timekit)

## 
## Attaching package: 'timekit'

## The following objects are masked from 'package:timetk':
## 
##     tk_augment_timeseries_signature, tk_get_timeseries_signature,
##     tk_get_timeseries_summary, tk_get_timeseries_unit_frequency,
##     tk_get_timeseries_variables, tk_index, tk_make_future_timeseries,
##     tk_tbl, tk_ts, tk_ts_, tk_ts_.data.frame, tk_ts_.default, tk_xts,
##     tk_xts_, tk_zoo, tk_zoo_, tk_zooreg, tk_zooreg_,
##     tk_zooreg_.data.frame, tk_zooreg_.default

```

```
library(broom)
library(ggthemes)
```

exploration

```
data %>% glimpse()

## # Rows: 41,607
## # Columns: 2
## $ time    <dttm> 2019-01-01 00:00:00, 2019-01-01 01:00:00, 2019-01-01 02:00:00~
## $ measure <dbl> 13.996154, 13.470886, 10.388608, 8.771795, 7.373077, 7.323377,~

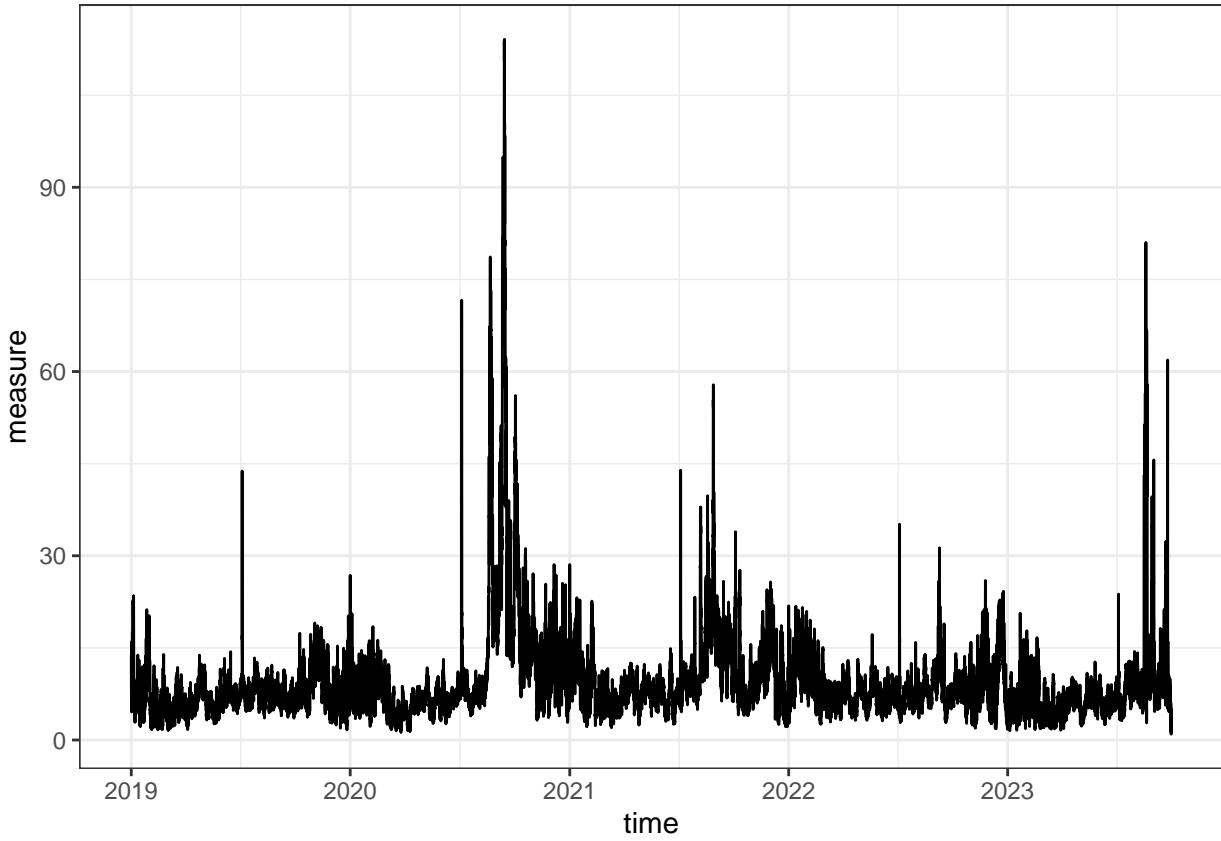
data %>%
  summarize_all(.funs = ~sum(is.na(.)))

## # A tibble: 1 x 2
##       time   measure
##     <int>     <int>
## 1      0         0

data %>%
  summarize(mean = mean(measure),
            max = max(measure),
            min = min(measure),
            IQR = IQR(measure))

## # A tibble: 1 x 4
##       mean   max   min   IQR
##     <dbl> <dbl> <dbl> <dbl>
## 1  9.19  114.  9.19  4.47

data %>%
  ggplot(aes(x = time, y = measure)) +
  geom_line() +
  theme_bw()
```



```
## train/test split
```

ts object

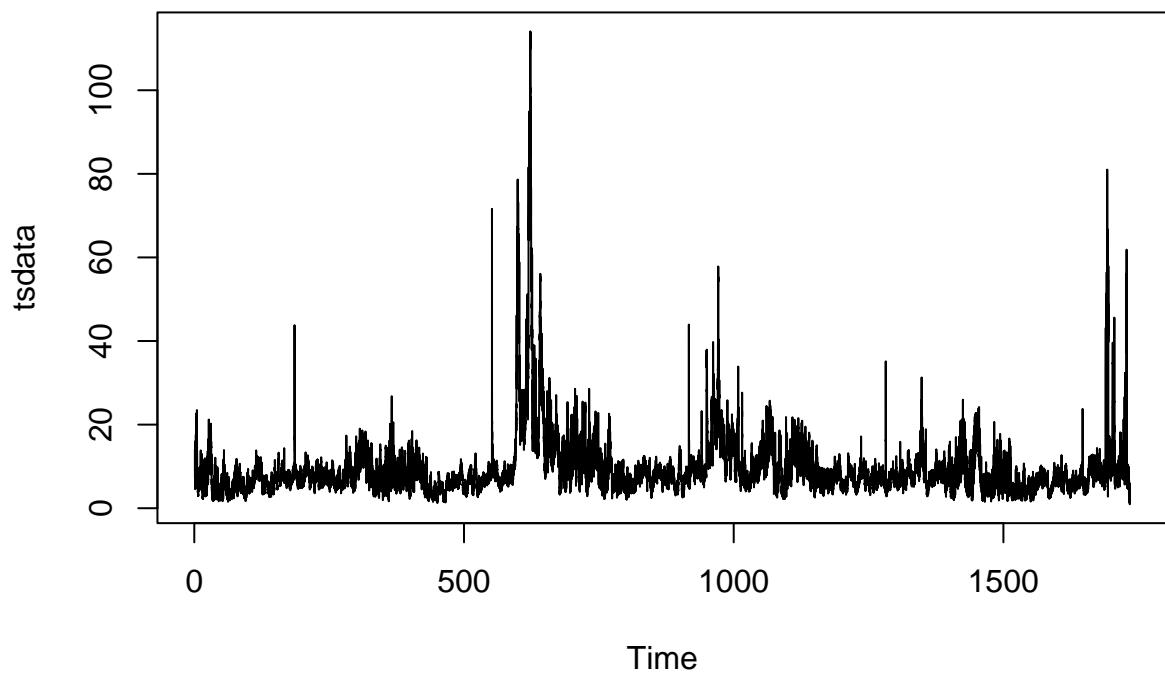
frequency is used to precise the frequency of data (here 24 is used to precise that is hourly data)

1. create the ts for all the data

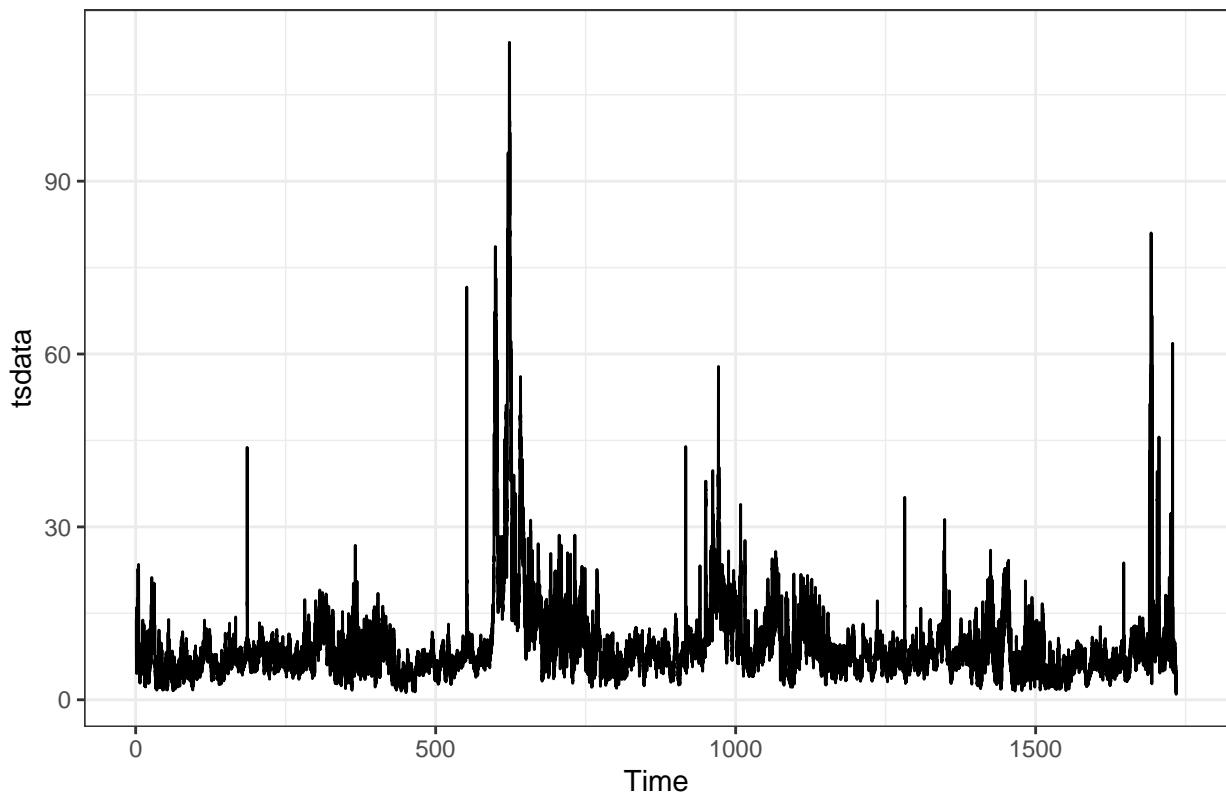
```
tsdata <- ts(data$measure,frequency = 24)
tsdata %>% head()
```

```
## Time Series:
## Start = c(1, 1)
## End = c(1, 6)
## Frequency = 24
## [1] 13.996154 13.470886 10.388608 8.771795 7.373077 7.323377
```

```
plot.ts(tsdata)
```



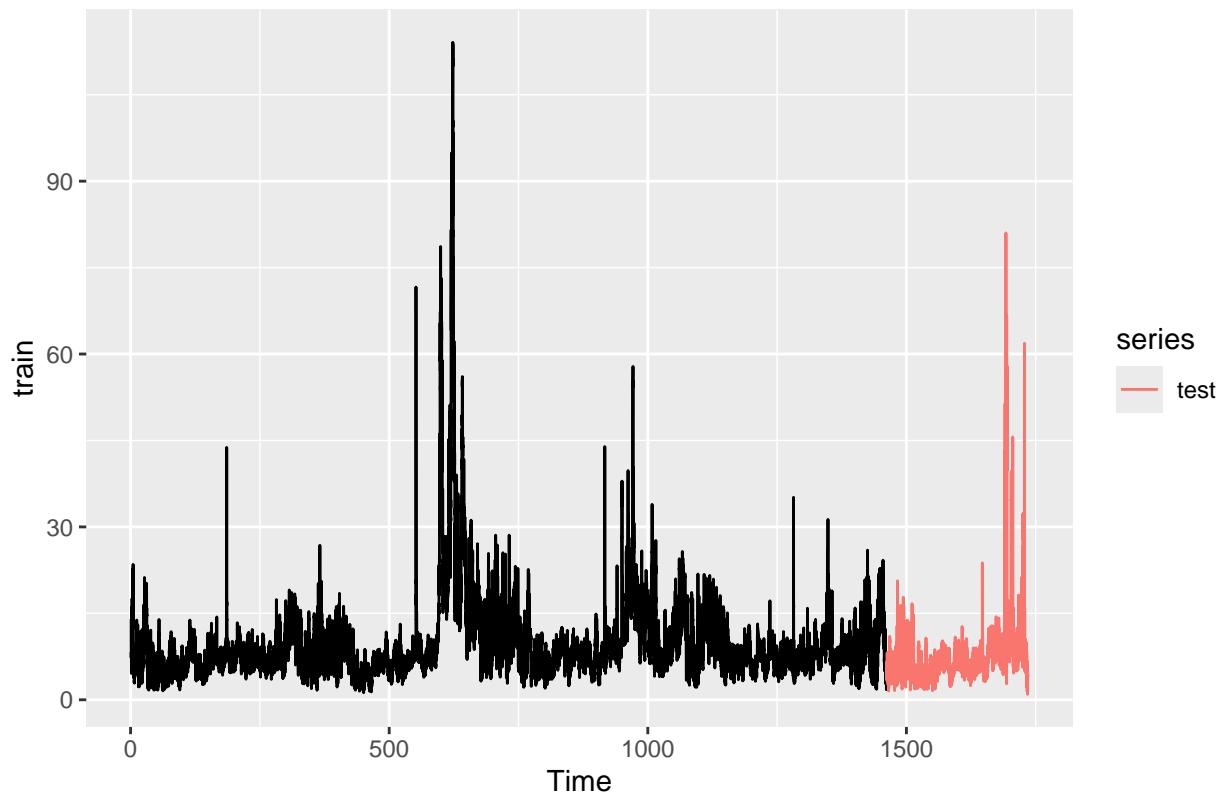
```
autoplot(tsdata) + theme_bw()
```



2. train 4 years : test = 2023

```
train <- tsdata %>% window(end = c(1461,24))
test <- tsdata %>% window(start = c(1462,1))
```

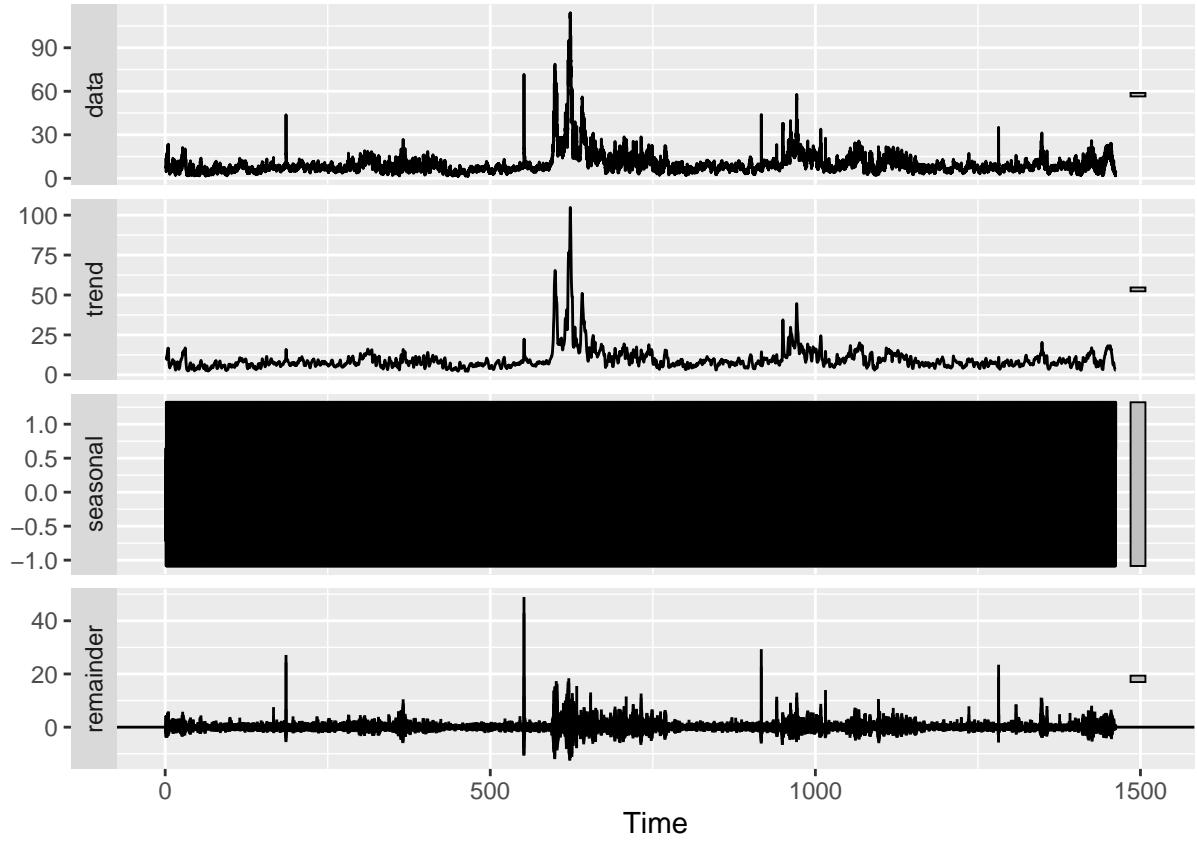
```
autoplot(train) +
  autolayer(test,series = "test")
```



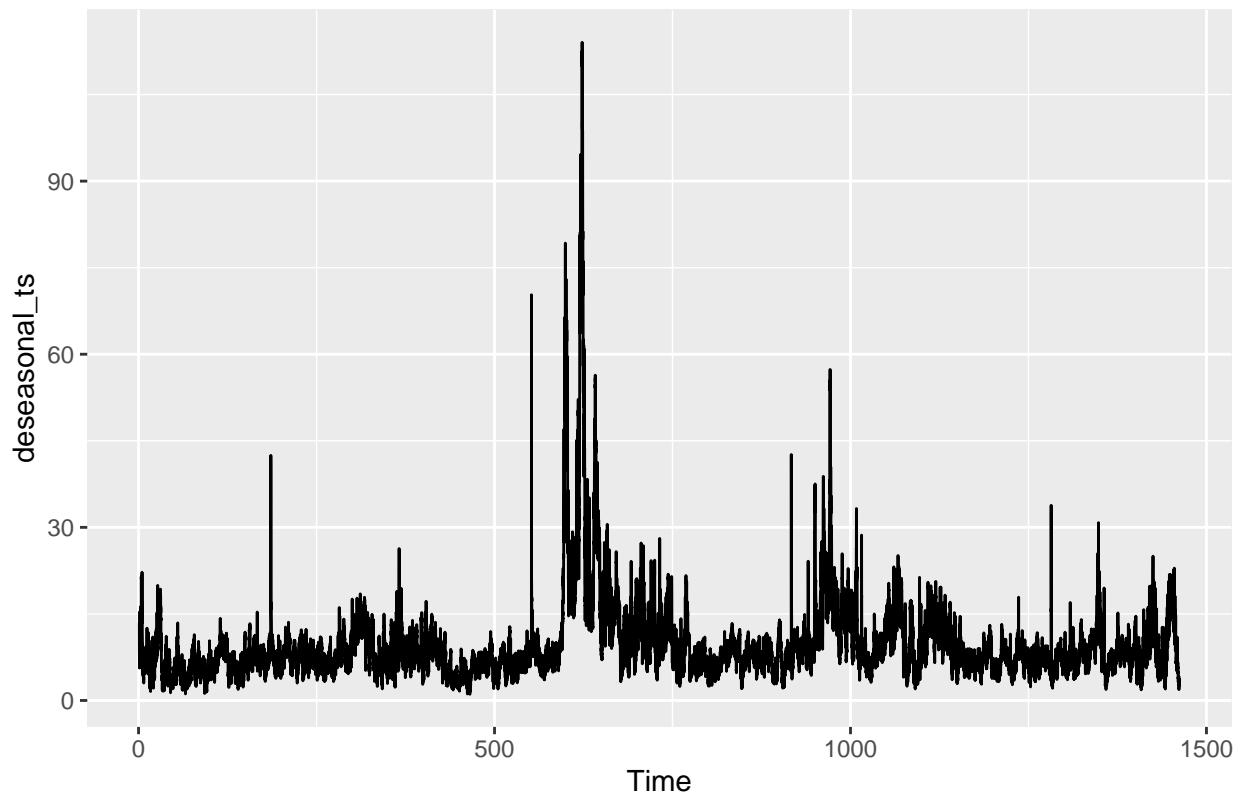
ARIMA model

stl is better than decompose() because it use non-linear approximation

```
decomp = stl(train,s.window = "periodic")
deseasonal_ts <- seasadj(decomp)
autoplot(decomp)
```



```
autoplot(deseasonal_ts)
```

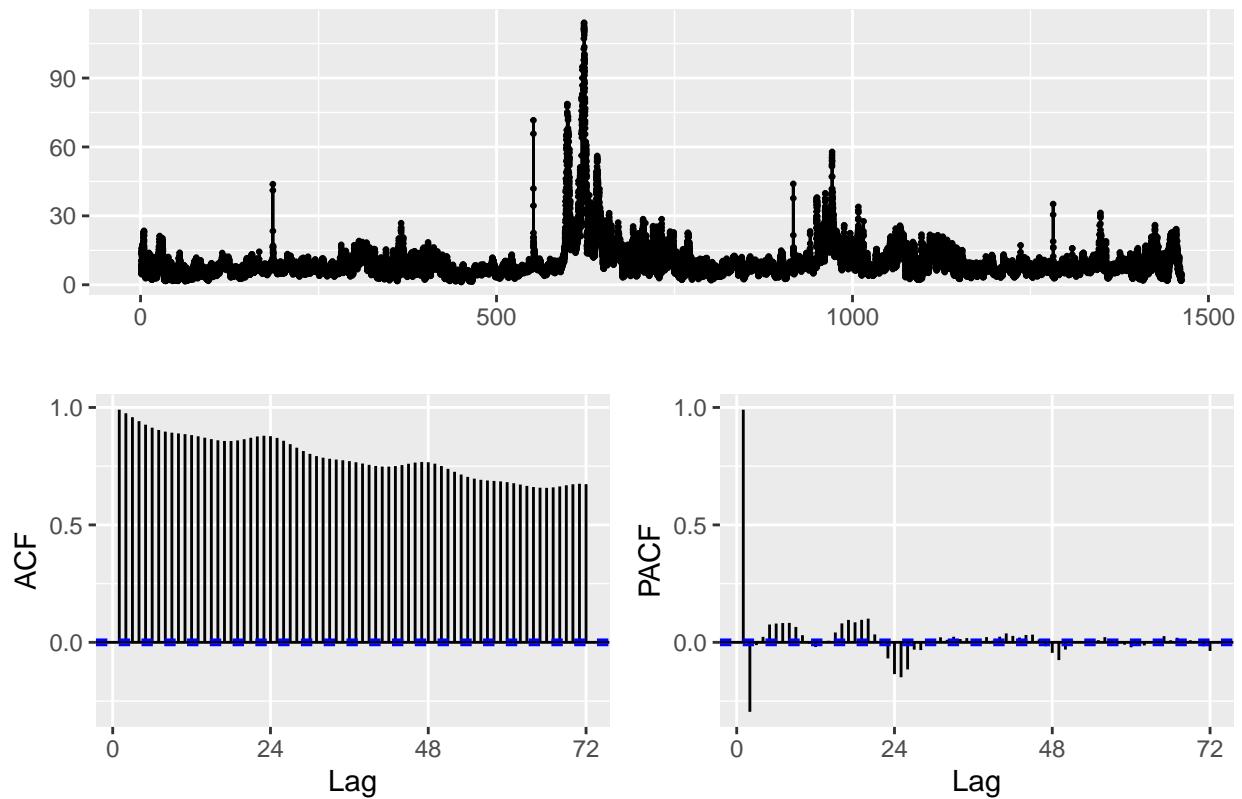


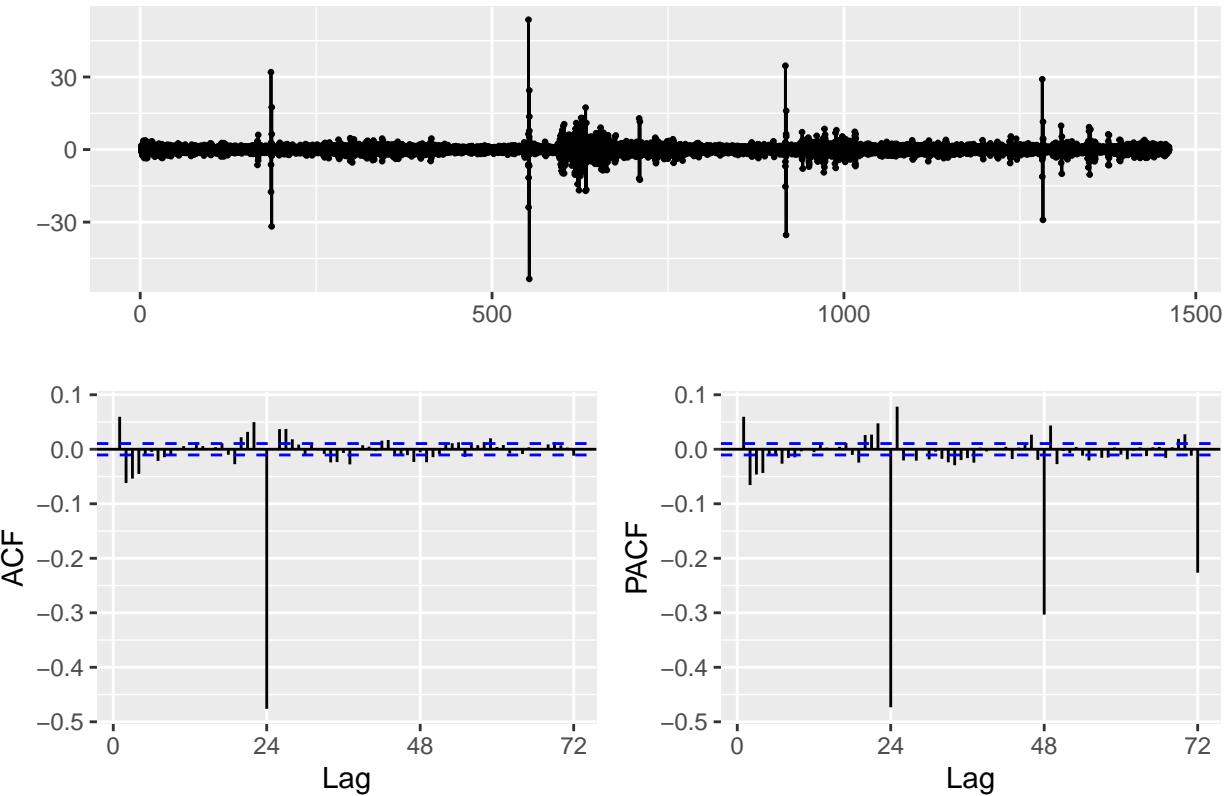
```
sweep::sw_tidy_decomp(decomp) %>% head()
```

```
## # A tibble: 6 x 6
##   index observed  season trend remainder seasadj
##   <dbl>    <dbl>    <dbl> <dbl>     <dbl>    <dbl>
## 1 1        14.0   0.486  9.19    4.32    13.5
## 2 1.04    13.5   0.0420  9.21    4.22    13.4
## 3 1.08    10.4   -0.360  9.23    1.52    10.7
## 4 1.12    8.77  -0.627  9.25    0.146   9.40
## 5 1.17    7.37  -0.718  9.27    -1.18   8.09
## 6 1.21    7.32  -0.586  9.30    -1.39   7.91
```

some ideas

```
ggtstdisplay(train)
```





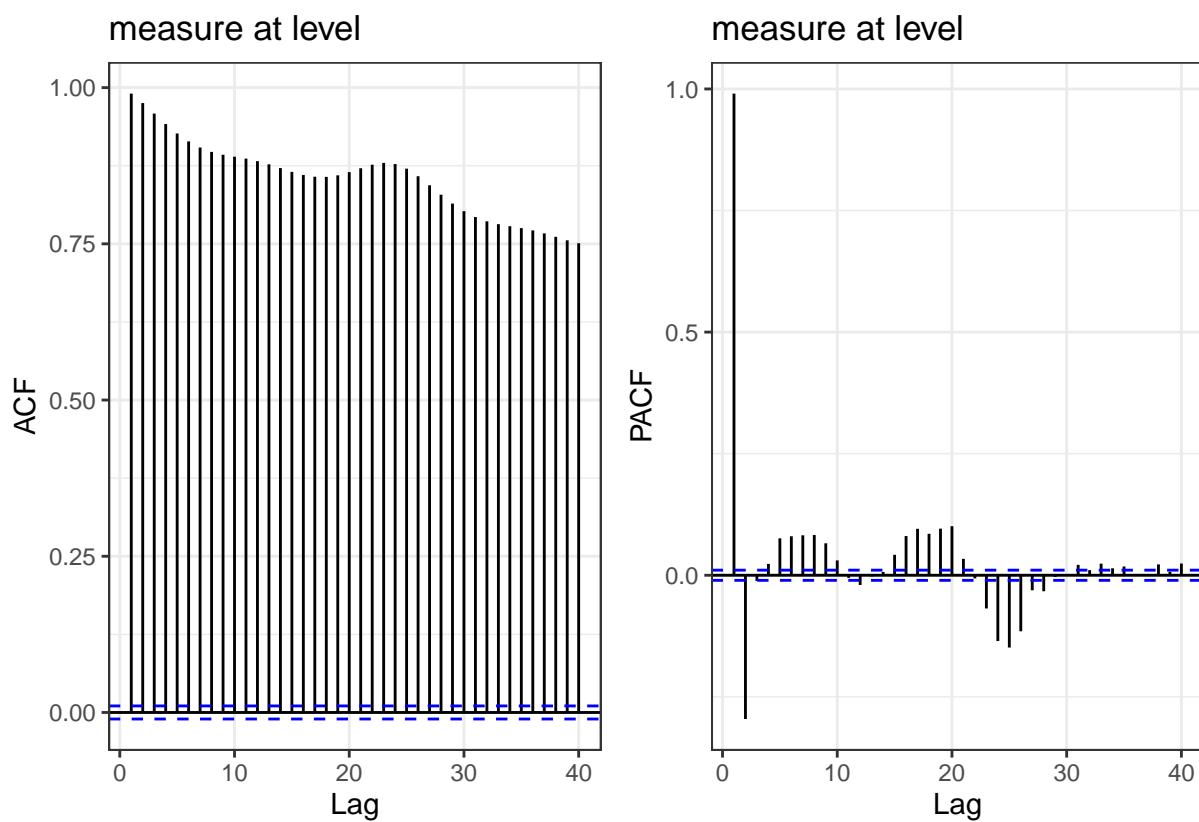
Autocorrelation plots

at level

```
library(patchwork)

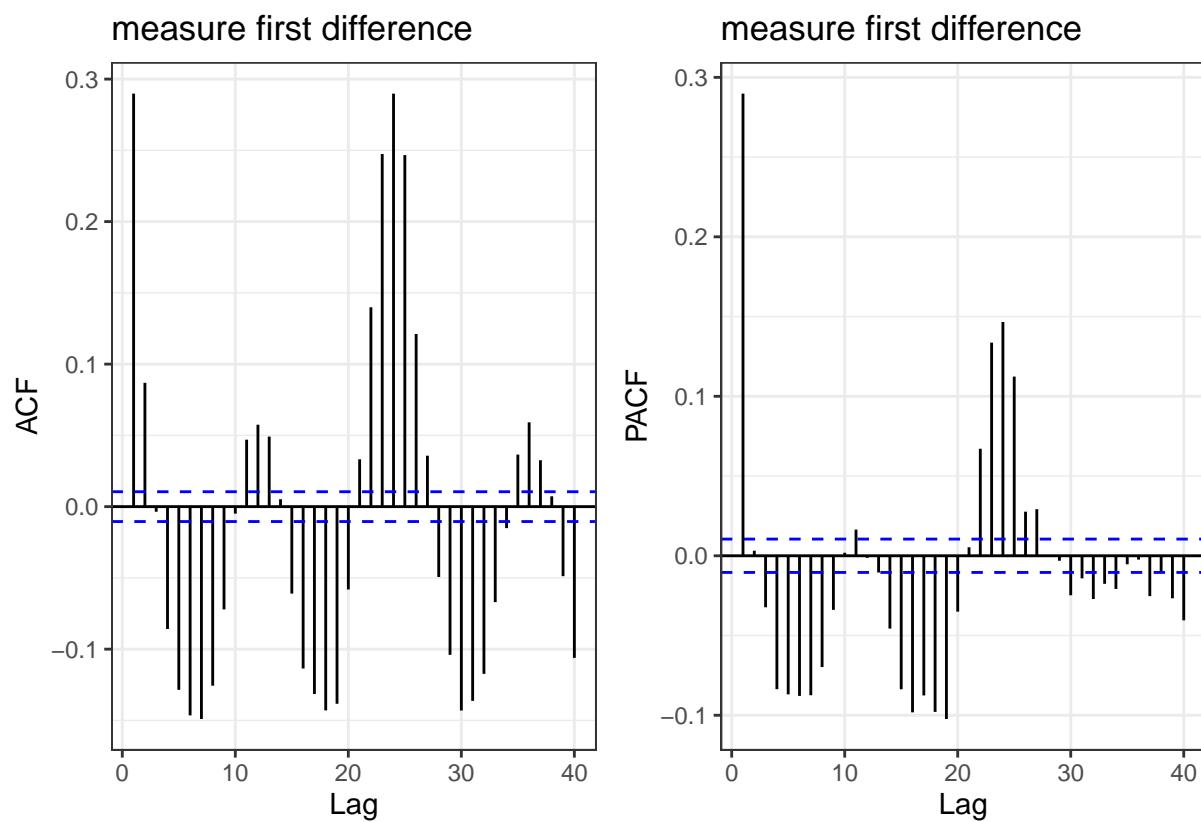
p1 <- ggAcf(train,lag=40) + labs(title="measure at level") + theme_bw()
p2 <- ggPacf(train,lag=40) + labs(title="measure at level") + theme_bw()

p1 + p2
```



```
p1 <- ggAcf(diff(train),lag=40) + labs(title="measure first difference") + theme_bw()
p2 <- ggPacf(diff(train),lag=40) + labs(title="measure first difference") + theme_bw()

p1 + p2
```

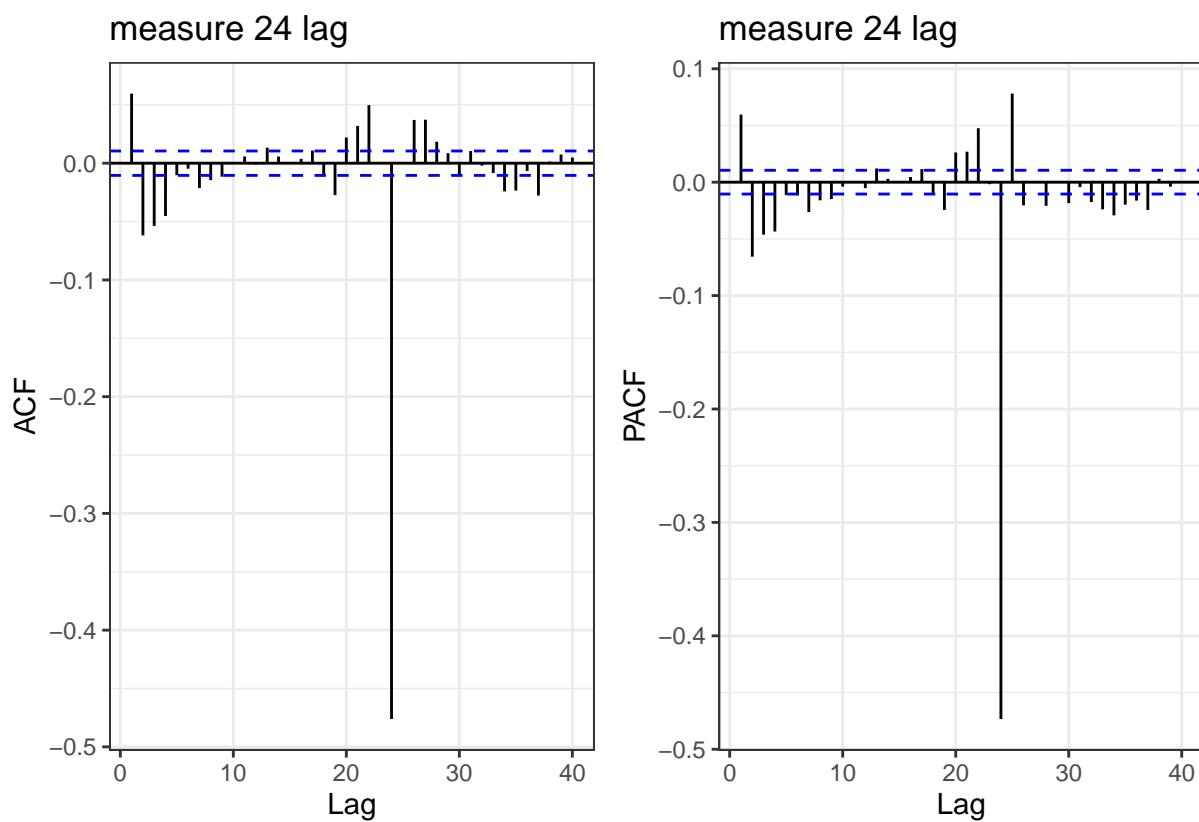


```

p1 <- ggAcf(diff(diff(train), lag = 24),lag=40) + labs(title="measure 24 lag") + theme_bw()
p2 <- ggPacf(diff(diff(train),lag = 24),lag=40) + labs(title="measure 24 lag") + theme_bw()

p1 + p2

```

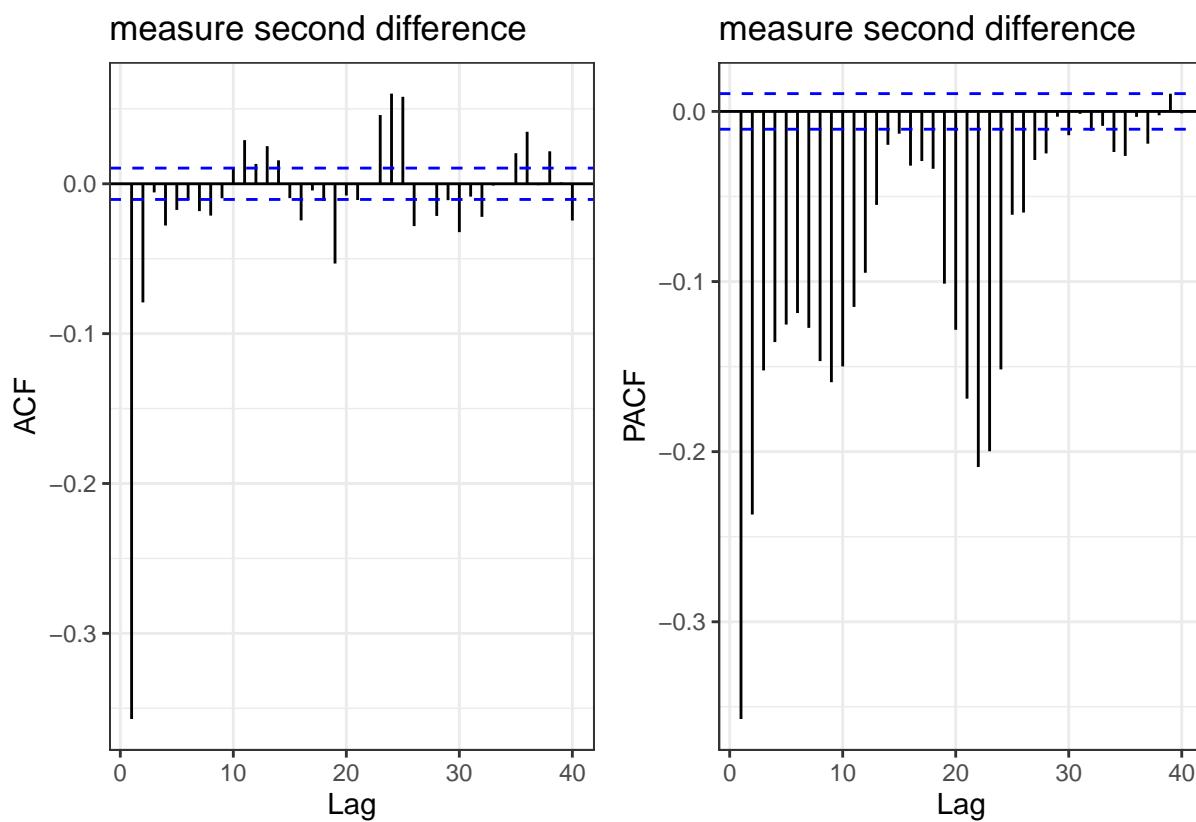


```

p1 <- ggAcf(diff(diff(train)),lag=40) + labs(title="measure second difference ") + theme_bw()
p2 <- ggPacf(diff(diff(train)),lag=40) + labs(title="measure second difference") + theme_bw()

p1 + p2

```



ADF test

level

```
sw_glance(adf.test(train))
```

```
## Warning in adf.test(train): p-value smaller than printed p-value
```

```
## # A tibble: 1 x 5
##   statistic p.value parameter method      alternative
##       <dbl>    <dbl>     <dbl> <chr>        <chr>
## 1     -10.5     0.01      32 Augmented Dickey-Fuller Test stationary
```

```
sw_glance(adf.test(diff(train,differences = 1)))
```

```
## Warning in adf.test(diff(train, differences = 1)): p-value smaller than printed
## p-value
```

```
## # A tibble: 1 x 5
##   statistic p.value parameter method      alternative
##       <dbl>    <dbl>     <dbl> <chr>        <chr>
## 1     -32.2     0.01      32 Augmented Dickey-Fuller Test stationary
```

```

sw_glance(adf.test(diff(train,differences = 2)))

## Warning in adf.test(diff(train, differences = 2)): p-value smaller than printed
## p-value

## # A tibble: 1 x 5
##   statistic p.value parameter method      alternative
##       <dbl>    <dbl>     <dbl> <chr>          <chr>
## 1      -59.4     0.01      32 Augmented Dickey-Fuller Test stationary

sw_glance(adf.test(diff(train,lag = 24)))

## Warning in adf.test(diff(train, lag = 24)): p-value smaller than printed
## p-value

## # A tibble: 1 x 5
##   statistic p.value parameter method      alternative
##       <dbl>    <dbl>     <dbl> <chr>          <chr>
## 1      -27.1     0.01      32 Augmented Dickey-Fuller Test stationary

sw_glance(adf.test(diff(train,differences = 1,lag = 24)))

## Warning in adf.test(diff(train, differences = 1, lag = 24)): p-value smaller
## than printed p-value

## # A tibble: 1 x 5
##   statistic p.value parameter method      alternative
##       <dbl>    <dbl>     <dbl> <chr>          <chr>
## 1      -27.1     0.01      32 Augmented Dickey-Fuller Test stationary

sw_glance(adf.test(diff(train,differences = 2,lag = 24)))

## Warning in adf.test(diff(train, differences = 2, lag = 24)): p-value smaller
## than printed p-value

## # A tibble: 1 x 5
##   statistic p.value parameter method      alternative
##       <dbl>    <dbl>     <dbl> <chr>          <chr>
## 1      -35.5     0.01      32 Augmented Dickey-Fuller Test stationary

```

Auto Arima

```
aar1 <- auto.arima(train, seasonal=TRUE,stepwise = T,trace = T,D = 1,max.p = 3,max.q = 3)
```

```

## 
## Fitting models using approximations to speed things up...
##
## ARIMA(2,0,2)(1,1,1)[24] with drift      : Inf
## ARIMA(0,0,0)(0,1,0)[24] with drift      : 190972.9
## ARIMA(1,0,0)(1,1,0)[24] with drift      : 104479.8
## ARIMA(0,0,1)(0,1,1)[24] with drift      : 155621.2
## ARIMA(0,0,0)(0,1,0)[24]                  : 190970.9
## ARIMA(1,0,0)(0,1,0)[24] with drift      : 113135.3
## ARIMA(1,0,0)(2,1,0)[24] with drift      : 101303.6
## ARIMA(1,0,0)(2,1,1)[24] with drift      : Inf
## ARIMA(1,0,0)(1,1,1)[24] with drift      : Inf
## ARIMA(0,0,0)(2,1,0)[24] with drift      : 190723.7
## ARIMA(2,0,0)(2,1,0)[24] with drift      : 100790.8
## ARIMA(2,0,0)(1,1,0)[24] with drift      : 104018.5
## ARIMA(2,0,0)(2,1,1)[24] with drift      : Inf
## ARIMA(2,0,0)(1,1,1)[24] with drift      : Inf
## ARIMA(3,0,0)(2,1,0)[24] with drift      : 100750.7
## ARIMA(3,0,0)(1,1,0)[24] with drift      : 103984.2
## ARIMA(3,0,0)(2,1,1)[24] with drift      : Inf
## ARIMA(3,0,0)(1,1,1)[24] with drift      : Inf
## ARIMA(3,0,1)(2,1,0)[24] with drift      : 100706.6
## ARIMA(3,0,1)(1,1,0)[24] with drift      : 103964.8
## ARIMA(3,0,1)(2,1,1)[24] with drift      : Inf
## ARIMA(3,0,1)(1,1,1)[24] with drift      : Inf
## ARIMA(2,0,1)(2,1,0)[24] with drift      : 100757.1
## ARIMA(3,0,2)(2,1,0)[24] with drift      : 100762
## ARIMA(2,0,2)(2,1,0)[24] with drift      : Inf
## ARIMA(3,0,1)(2,1,0)[24]                  : 100704.6
## ARIMA(3,0,1)(1,1,0)[24]                  : 103962.8
## ARIMA(3,0,1)(2,1,1)[24]                  : Inf
## ARIMA(3,0,1)(1,1,1)[24]                  : Inf
## ARIMA(2,0,1)(2,1,0)[24]                  : 100755.1
## ARIMA(3,0,0)(2,1,0)[24]                  : 100748.7
## ARIMA(3,0,2)(2,1,0)[24]                  : 100760
## ARIMA(2,0,0)(2,1,0)[24]                  : 100788.8
## ARIMA(2,0,2)(2,1,0)[24]                  : Inf
##
## Now re-fitting the best model(s) without approximations...
##
## ARIMA(3,0,1)(2,1,0)[24]                  : Inf
## ARIMA(3,0,1)(2,1,0)[24] with drift      : Inf
## ARIMA(3,0,0)(2,1,0)[24]                  : Inf
## ARIMA(3,0,0)(2,1,0)[24] with drift      : Inf
## ARIMA(2,0,1)(2,1,0)[24]                  : Inf
## ARIMA(2,0,1)(2,1,0)[24] with drift      : Inf
## ARIMA(3,0,2)(2,1,0)[24]                  : Inf
## ARIMA(3,0,2)(2,1,0)[24] with drift      : Inf
## ARIMA(2,0,0)(2,1,0)[24]                  : Inf
## ARIMA(2,0,0)(2,1,0)[24] with drift      : Inf
## ARIMA(1,0,0)(2,1,0)[24] with drift      : 101358.4
##
## Best model: ARIMA(1,0,0)(2,1,0)[24] with drift

```

```
sw_tidy(aar1)
```

```
## # A tibble: 4 x 2
##   term    estimate
##   <chr>     <dbl>
## 1 ar1      0.971
## 2 sar1     -0.613
## 3 sar2     -0.296
## 4 drift    -0.000182
```

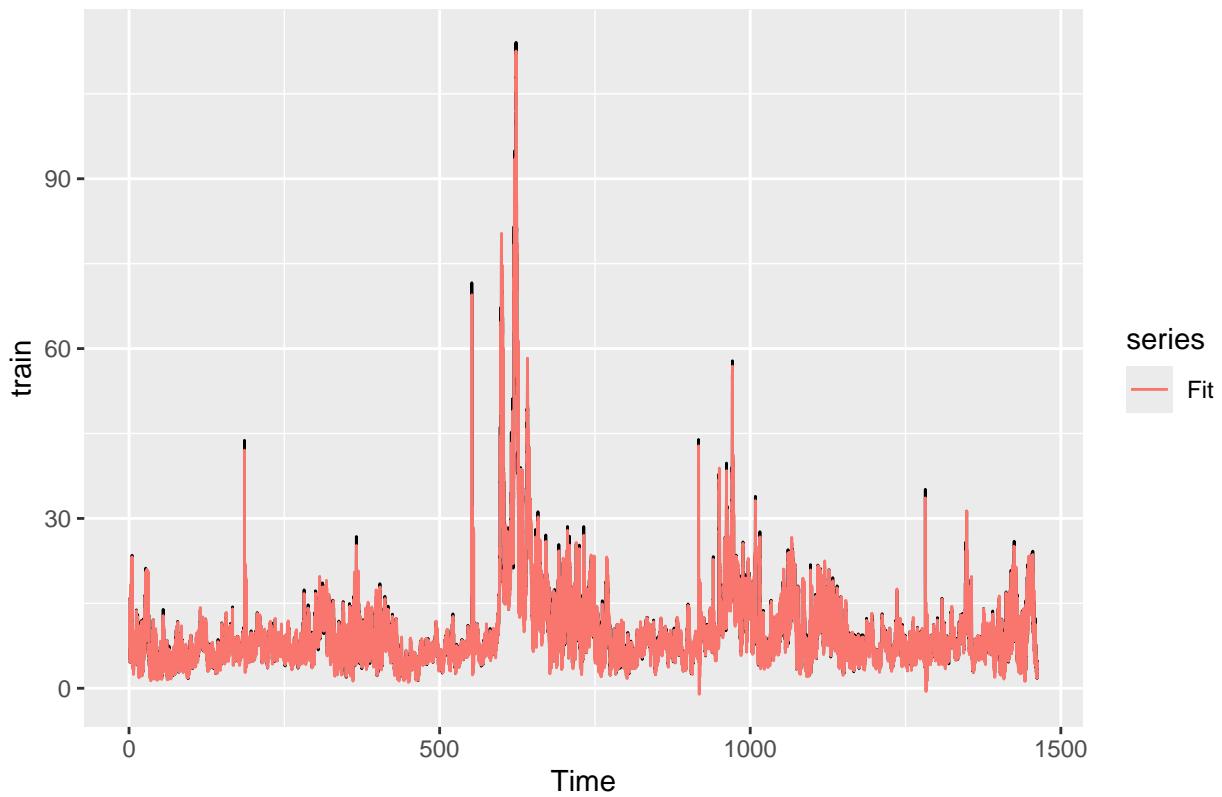
```
sw_glance(aar1)
```

```
## # A tibble: 1 x 12
##   model.desc sigma logLik   AIC     BIC       ME   RMSE    MAE    MPE    MAPE    MASE
##   <chr>     <dbl>  <dbl>  <dbl>  <dbl>  <dbl> <dbl>  <dbl>  <dbl> <dbl>  <dbl>
## 1 ARIMA(1,0~  1.03 -50674. 1.01e5 1.01e5 -3.56e-5 1.03 0.570 -0.358 6.61 0.277
## # i 1 more variable: ACF1 <dbl>
```

```
sw_augment(aar1) %>% head()
```

```
## # A tibble: 6 x 4
##   index .actual .fitted .resid
##   <dbl>   <dbl>   <dbl>   <dbl>
## 1 1      14.0    14.0   0.0140
## 2 1.04   13.5    13.5   0.0135
## 3 1.08   10.4    10.4   0.0104
## 4 1.12   8.77    8.76   0.00877
## 5 1.17   7.37    7.37   0.00737
## 6 1.21   7.32    7.32   0.00732
```

```
autoplot(train) +
  autolayer(aar1$fitted, series="Fit")
```



forecast on test

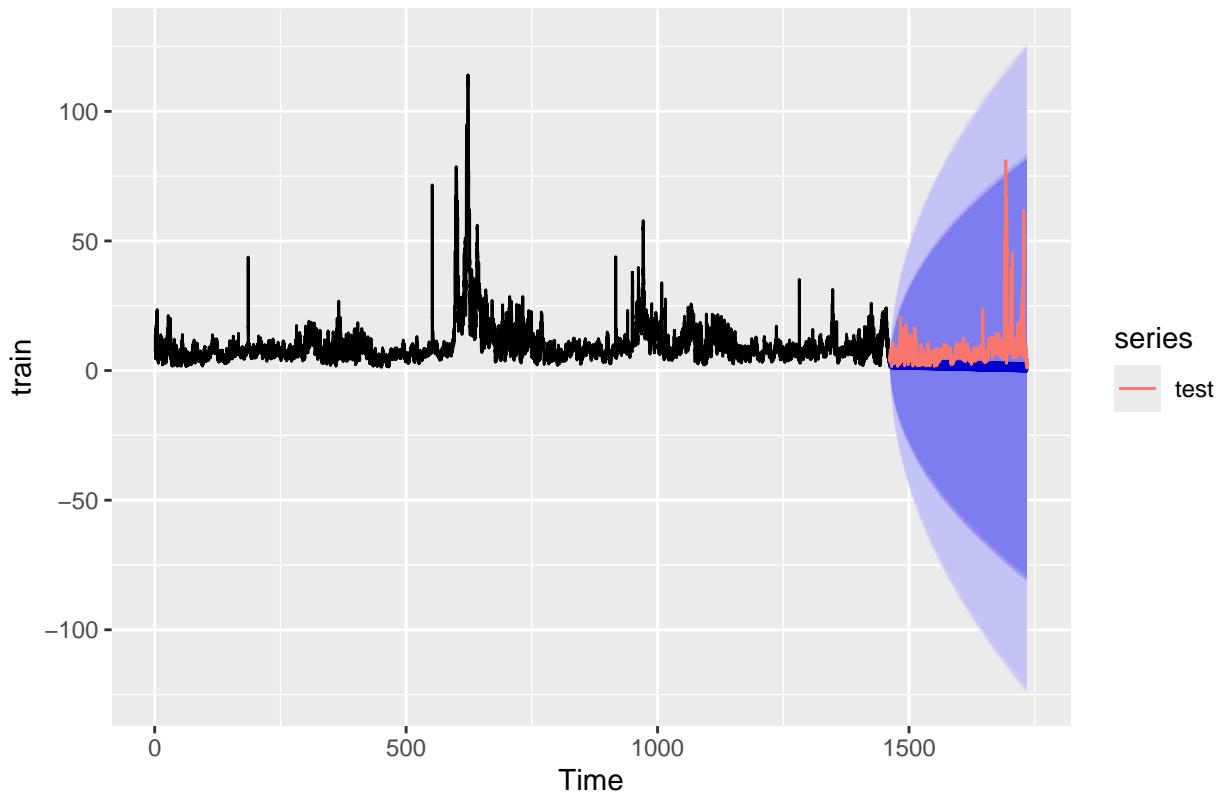
```
f1 <- forecast(aar1,h = 6543)
#f1 %>% head()
```

```
summary(aar1)
```

```
## Series: train
## ARIMA(1,0,0)(2,1,0)[24] with drift
##
## Coefficients:
##             ar1      sar1      sar2     drift
##            0.9709   -0.6127   -0.2958   -0.0002
## s.e.    0.0013    0.0051    0.0051    0.0041
##
## sigma^2 = 1.056: log likelihood = -50674.19
## AIC=101358.4   AICc=101358.4   BIC=101400.7
##
## Training set error measures:
##                  ME      RMSE       MAE       MPE      MAPE      MASE
## Training set -3.557011e-05 1.027079 0.5700497 -0.3576875 6.614659 0.2771442
##                  ACF1
## Training set 0.1166776
```

```
aar1 %>%
  forecast(h=6543) %>%
  autoplot() + autolayer(test)
```

Forecasts from ARIMA(1,0,0)(2,1,0)[24] with drift



```
accuracy(f1 , test)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -3.557011e-05 1.027079 0.5700497 -0.3576875 6.614659 0.2771442
## Test set      6.083488e+00 8.373799 6.0897210 73.5912045 73.857791 2.9606734
##                  ACF1 Theil's U
## Training set 0.1166776       NA
## Test set     0.9426066  3.062574
```

comments :

no the best of results needs some reflexion !!